

---

## COS 350: Program #6: multiple file word counting

Objectives: Working with fork, exec, pipes and comparing with pthreads and shared memory communication.

---

You will write 3 versions of a multiple file word counting program. Your goal is to make all three versions as fast as possible. Do not use `getc()` as in the class examples. Use `read()` to read in a large chunk of the file into a buffer for processing. This will be faster.

### Version 1: mywc1

Works like `wc` (word count) but simply counts the number of words in each file specified on the command line, and for each file it prints the count and the filename, and finally the total when there is more than 1 file. Your results should match `wc -w`.

A word is considered to be any non-zero length sequence of characters delimited by white space. A white space character is a space, tab, newline or other such character as identified by the `isspace(c)` function from `ctype.h`.

### Version 2: mywc2

Use `fork()` and `exec()` to create a separate process for each individual file. Just `exec()` `mywc1` to do the counting and setup pipes to read back the output so that the main program can total them up. The output should be identical to `mywc1` except that the output order is allowed to vary. Fork them all off so that they are working simultaneously. This should be faster than `mywc1` when run on multiple files.

### Version 3: mywc3

Use `pthread`s to count the files. Create simultaneous threads, one for each file. You can decide what sort of synchronization you need. As with `mywc2`, the output order may vary. Remember to use the `-pthread` compiler flag. Again, this should be faster than `mywc1` when run on multiple files.

### Testing:

Once you have your programs all working, recompile them using the optimization flag `-O2` (e.g. `gcc -O2 ...`). This will make them run faster.

Use the unix program `time` for timing your programs. Time each of these running on the set of 10 files: `f1 – f10` (these are in the course directory). Ideally, do the timing tests on an unloaded machine. You can check the load on a machine with the program: `uptime`. The load values should be  $< 1.0$ . Preferably close to 0.0.

Run at least 3 trials for each version and report the best time among these trials. Which version is fastest? Try to guess why.

## What to turn in:

### Written report:

1. Your name(s) and if a team, who did the electronic submit.
2. A discussion of any incomplete parts, known bugs, deviations from the specification, or extra features in your program.
3. Your code
4. Your Makefile
5. A log of the following testing results. (I recommend using script to capture it.)

```
time mywc1 f*           test mywc1 counting all 10 files f1-f10
time mywc2 f*           test mywc2
time mywc3 f*           test mywc3
exit                    exit script
```

6. A neatly organized table of your timing results like this:

Machine that the timing tests were performed on: \_\_\_\_\_

Try to use an unloaded machine and report the best of at least 3 trials for each program.

Timings results	real	user	sys
mywc1			
mywc2			
mywc3			

7. Discuss which version was the fastest, and why you think that was?

### Electronic submission:

- Before you submit your program, compile and test it on a linux machine in the lab.
- From a machine in the lab run the program “submit” to submit your files.
- Submit your source code (mywc1.c mywc2.c mywc3.c) and any other source files, a Makefile that works for all 3 files, and your executables (mywc1 mywc2 mywc3) to a directory named: prog6
- Do not in any way combine, compress, zip, or tar your files!

### Grading:

- 5 points: Makefile
- 10 points: mywc1 gets correct individual file word counts
- 5 points: mywc1 gets correct total
- 15 points: mywc2 the processes execute concurrently, and it is faster than mywc1
- 15 points: mywc2 gets correct total
- 15 points: mywc3 the threads execute concurrently, and it is faster than mywc1
- 15 points: mywc3 gets correct total
- 10 points: Timing results and discussion
- 10 points: all versions finish in under 1 second